



Greedy Algorithms

Yan Gu



Greedy Algorithms

- Among the commonly-used algorithm design strategies, greedy probably is the most intuitive and easiest to understand
- Once decision at a time
- When you need to make a decision, choose the “best” based on a certain criterion
- Not necessarily optimal, need to prove it

What do greedy choice and optimal substructure mean?

- **Greedy choice (intuitively):**
 - The element t you greedily choose is not a bad idea!
 - It appears in some optimal solution!
 - (for any optimal solution, if it doesn't contain t , we can modify it to contain t !)
 - So just choose it!
- **Optimal substructure (intuitively):**
 - After choosing some element t
 - The final optimal solution is just to find the optimal solution for the rest of the (compatible) elements!
 - Recursively solve it using the same approach
- **So we repeatedly choose the greedy choice!**

Well, sometimes greedy is not optimal, is it useless?

- They may still provide you with some nice features!

Approximation Algorithms and Competitive Analysis

Ski Rental (rent or buy problem)

- You recently came to UCR and realized that you can go skiing at the Big Bear Lake Resort
- Buying the equipment costs about \$500 and renting it for a weekend costs \$50. Should you buy or rent?
- Clearly it depends on how many more times you go skiing in the future
 - If you will go skiing a total of 11 times or more, then it is better to buy, and to do it now
 - If we will go 9 times or fewer, then it is better to rent
 - If we go 10 times it does not matter



A reasonable strategy

- You will rent it first and buy it later
- You will first rent it for the first 10 times, and buy it if you go for the 11th time
 - If you go 10 times or fewer time, then the strategy is optimal
 - If you go for the 11th time, you pay $10 \times \$50 + \$500 = \$1000$, while the optimal strategy (buying it on the first time) pays \$500
 - You will never pay more than twice of the optimal strategy



Approximation algorithms


- For a minimization problems instance I and an algorithm ALG , let $ALG(I)$ be the quality of the algorithm's output and $OPT(I)$ be the quality of the optimal solution
- For $c > 1$, ALG is a c -approximation algorithm if for every I , $ALG(I) \leq c \cdot OPT(I)$
- The abovementioned is a 2-approximation algorithm for the (online) ski rental problem

Approximation algorithms

- For a maximization problems instance I and an algorithm ALG , let $ALG(I)$ be the quality of the algorithm's output and $OPT(I)$ be the quality of the optimal solution
- For $c < 1$, ALG is a c -approximation algorithm if for every I , $ALG(I) \geq c \cdot OPT(I)$


Holiday Celebration!

- You work for Riverside City Council and arranging the next Christmas celebration!
- You have the map of the celebration region, and want to have some police on some intersections, so they can help arranging and provide assistance
- You want to minimize the “police stops” but still have police on one of the two intersections of every (segment) of street
- Unfortunately, solving this problem (vertex cover) on a general graph is NP-Complete




Vertex Cover Problem

- It is hard to know the exact number M , but ...
- Consider a simple strategy, I enumerate every street, and if it is “uncovered”, I just put police stops on both ends, and let S be the stops we put here
- Lemma: $|S| \geq |M|$
- Proof: S must cover at least one vertex for each edge in M , and this vertex covers no other edge in M



An approximate solution

- Repeatedly find “uncovered” streets and put stops on both ends
- Lemma: $|S| \leq 2 \cdot OPT(G)$
- Proof: For each $e \in E$, at least one endpoint is in M , so S is a valid solution, and $|S| = 2|M| \leq 2 \cdot OPT(G)$
- No need to know $OPT(G)$!



Fun fact: can we replace the constant 2 with $\alpha < 2$?

- No better constant-factor approximation algorithm is known
- The minimum vertex cover problem is APX-complete: it cannot be approximated arbitrarily well unless $P = NP$
- Using PCP theorem, one can show that $\alpha \geq \sqrt{2}$; if the unique games conjecture is true, then $\alpha = 2$
- More similar ideas will be covered by CS 219 by Amey
- Another course: <https://www.cs.cmu.edu/~anupamg/adv-approx/>


Bin Packing

- Given n items of different weights and bins each of capacity c , assign each item to a bin such that number of total used bins is minimized. It may be assumed that all items have weights smaller than bin capacity.

Bin Packing


- Given n items of different weights and bins each of capacity c , assign each item to a bin such that number of total used bins is minimized. It may be assumed that all items have weights smaller than bin capacity.

Example: Bin Packing




Using the fewest number of bins to pack all the objects

Example: Bin Packing




Using the fewest number of bins to pack all the objects

Example: Bin Packing




Using the fewest number of bins to pack all the objects

Example: Bin Packing




Using the fewest number of bins to pack all the objects

Example: Bin Packing




Using the fewest number of bins to pack all the objects

Example: Bin Packing



Number of bins = 3

Example: Bin Packing



(Optimal Solution) Number of bins = 2

Bin Packing

- Given n items of different weights and bins each of capacity c , assign each item to a bin such that number of total used bins is minimized. It may be assumed that all items have weights smaller than bin capacity.
- What strategies will you use?


Greedy algorithms and their bounds


- If the optimal solution is x
- Next fit (put in the current, start a new one if it doesn't fit the current): no more than $2x$ bins
- First fit (find the first one that fits): no more than $1.7x$ bins
- Best fit (put in the tightest spot): no more than $1.7x$ bins
- Worst fit (leave the largest space): no more than $2x - 2$ bins
- First fit decreasing (sort and work on the largest first): no more than $(4x + 1)/3$ bins

Maximum Coverage Problem

The maximum coverage problem


- You will organize a seminar at UCR!
- You want to invite some famous researchers, each of them will present their recent papers
- One paper may have multiple authors – anyone of them can give the talk
- A researcher may have multiple papers – they can present all of them
- Well, but you only have budget to invite k of them
- You want to maximize the number of papers that may be presented!
- (The general question: given a bipartite graph, choose k vertices on the left to cover as many vertices on the right as possible)





$$\begin{aligned}n &= 5 \\m &= 8 \\k &= 3\end{aligned}$$

How to maximize the number of presented papers?




$$n = 5$$
$$m = 8$$
$$k = 3$$

How to maximize the number of presented papers?


$$n = 5$$

$$m = 8$$

$$k = 3$$



How to maximize the number of presented papers?



- If $k = 1$,
 - Choose the one with the most number of papers!
 - Either Alice or Bob
- If $k = 2$, should we choose the top two with #papers?
 - Not necessarily!
 - Alice + Bob = 5
 - Alice + David (or Erin) = 6
- What is the criteria that we should consider??

$$n = 5$$


$$m = 8$$

$$k = 3$$

How to maximize the number of presented papers?

Marginal Increase

- Given a set V , we want to select a subset $S \subseteq V$, in order to maximize some function $f(S)$
 - $f(S)$ is a function on a set
- When adding a new element x to the set, we say the “marginal gain” or “marginal increase” of adding x is
$$f(S \cup \{x\}) - f(S)$$
- Which is the extra benefit we can get by adding x
- For the researcher-paper problem, if we already decide to invite some researchers (set S), the marginal increase of inviting another researcher is the number of extra talks s/he will give



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice, who is the one with the most “marginal gain”?

How to maximize the number of presented papers?

Marginal
gain

1

Alice




Bob





Charlie



David



Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice, who is the one with the most “marginal gain”?

- $f(\{\text{Alice, Bob}\}) = 5$

How to maximize the
number of presented
papers?

Marginal
gain

1

1

Alice



Bob




Charlie



David



Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice, who is the one with the most “marginal gain”?
- $f(\{\text{Alice, Bob}\}) = 5$
- $f(\{\text{Alice, Charlie}\}) = 5$

How to maximize the number of presented papers?

Marginal
gain

1

Alice



Bob



1

Charlie




2

David



Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice, who is the one with the most “marginal gain”?
- $f(\{\text{Alice, Bob}\}) = 5$
- $f(\{\text{Alice, Charlie}\}) = 5$
- $f(\{\text{Alice, David}\}) = 6$

How to maximize the number of presented papers?


Marginal
gain

1

Alice



Bob



1

Charlie





2

David



2

Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice, who is the one with the most “marginal gain”?

- $f(\{\text{Alice, Bob}\}) = 5$
- $f(\{\text{Alice, Charlie}\}) = 5$
- $f(\{\text{Alice, David}\}) = 6$
- $f(\{\text{Alice, Erin}\}) = 6$

How to maximize the number of presented papers?

Marginal
gain

1

Alice



Bob



1

Charlie





2

David



2

Erin



$$n = 5$$


$$m = 8$$

$$k = 3$$

- $f(\{\text{Alice, Erin}\}) = 6$

How to maximize the
number of presented
papers?

Marginal
gain



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice and Erin, who is the one with the most “marginal gain”?

- $f(\{Alice, Erin\}) = 6$

How to maximize the
number of presented
papers?

Marginal
gain

1

Alice



Bob




Charlie



David



Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice and Erin, who is the one with the most “marginal gain”?

$$f(\{Alice, Erin, Bob\}) = 7$$

How to maximize the
number of presented
papers?

Marginal
gain

Alice




1 Bob





1 Charlie



David



Erin



$$n = 5$$

$$m = 8$$

$$k = 3$$

- If we have invited Alice and Erin, who is the one with the most “marginal gain”?

- $f(\{\text{Alice, Erin, Bob}\}) = 7$
- $f(\{\text{Alice, Erin, Charlie}\}) = 7$

How to maximize the
number of presented
papers?

Marginal
gain

1

1

0

Alice



Bob




Charlie



David



Erin



$$n = 5$$

$$m = 8$$


$$k = 3$$

- If we have invited Alice and Erin, who is the one with the most “marginal gain”?

- $f(\{\text{Alice, Erin, Bob}\}) = 7$
- $f(\{\text{Alice, Erin, Charlie}\}) = 7$
- $f(\{\text{Alice, Erin, David}\}) = 6$

How to maximize the
number of presented
papers?

Marginal
gain



$$n = 5$$

$$m = 8$$


$$k = 3$$

- If we have invited Alice and Erin, who is the one with the most “marginal gain”?

$$f(\{Alice, Erin, Bob\}) = 7$$

How to maximize the
number of presented
papers?

Marginal
gain



$$n = 5$$

$$m = 8$$


$$k = 3$$

Will “**marginal
gain first**” give
you the best
solution?

- $f(\{\text{Alice, Erin, Bob}\}) = 7$

How to maximize the
number of presented
papers?

Marginal
gain



$$n = 5$$
$$m = 8$$
$$k = 3$$

Will “**marginal
gain first**” give
you the best
solution?

Maximum Coverage Problem

- Using the “marginal gain first” strategy, the solution is
 - Not necessary optimal
 - But, can be **at least 63%** as good as the best solution!
 - I.e., the greedy algorithm gives you a 0.63-approximation
 - The greedy algorithm is essentially the **best-possible polynomial time approximation** algorithm for maximum coverage unless $P = NP$
 - This “63%” is actually $1 - \frac{1}{e}$

Submodularity

- When your objective is a **monotone submodular** function, this “marginal gain first” strategy will provide you a solution that is $\left(1 - \frac{1}{e}\right)$ -competitive

- Monotone: For sets S and $T \subseteq S$, $f(T) \leq f(S)$
- Submodular: For set S and $T \subseteq S$, and $x \notin S$,
$$f(T \cup \{x\}) - f(T) \geq f(S \cup \{x\}) - f(S)$$

=> adding an element to a smaller set will have a better marginal gain than adding it to a large set!

- (You can try to verify this on the maximum coverage problem)
- The greedy algorithm will find a solution that is at least $(1-1/e) \approx .63$ as good as optimal

Stable Matching Problem

Finding a good match

- n companies and n candidates
- Each candidate will go to one company, and each company will hire one candidate
- “preference list”

Alice		UGLF		ADCB
Bob		ULGF		ABCD
Charlie		FUGL		ACDB
David		LGUF		BADC

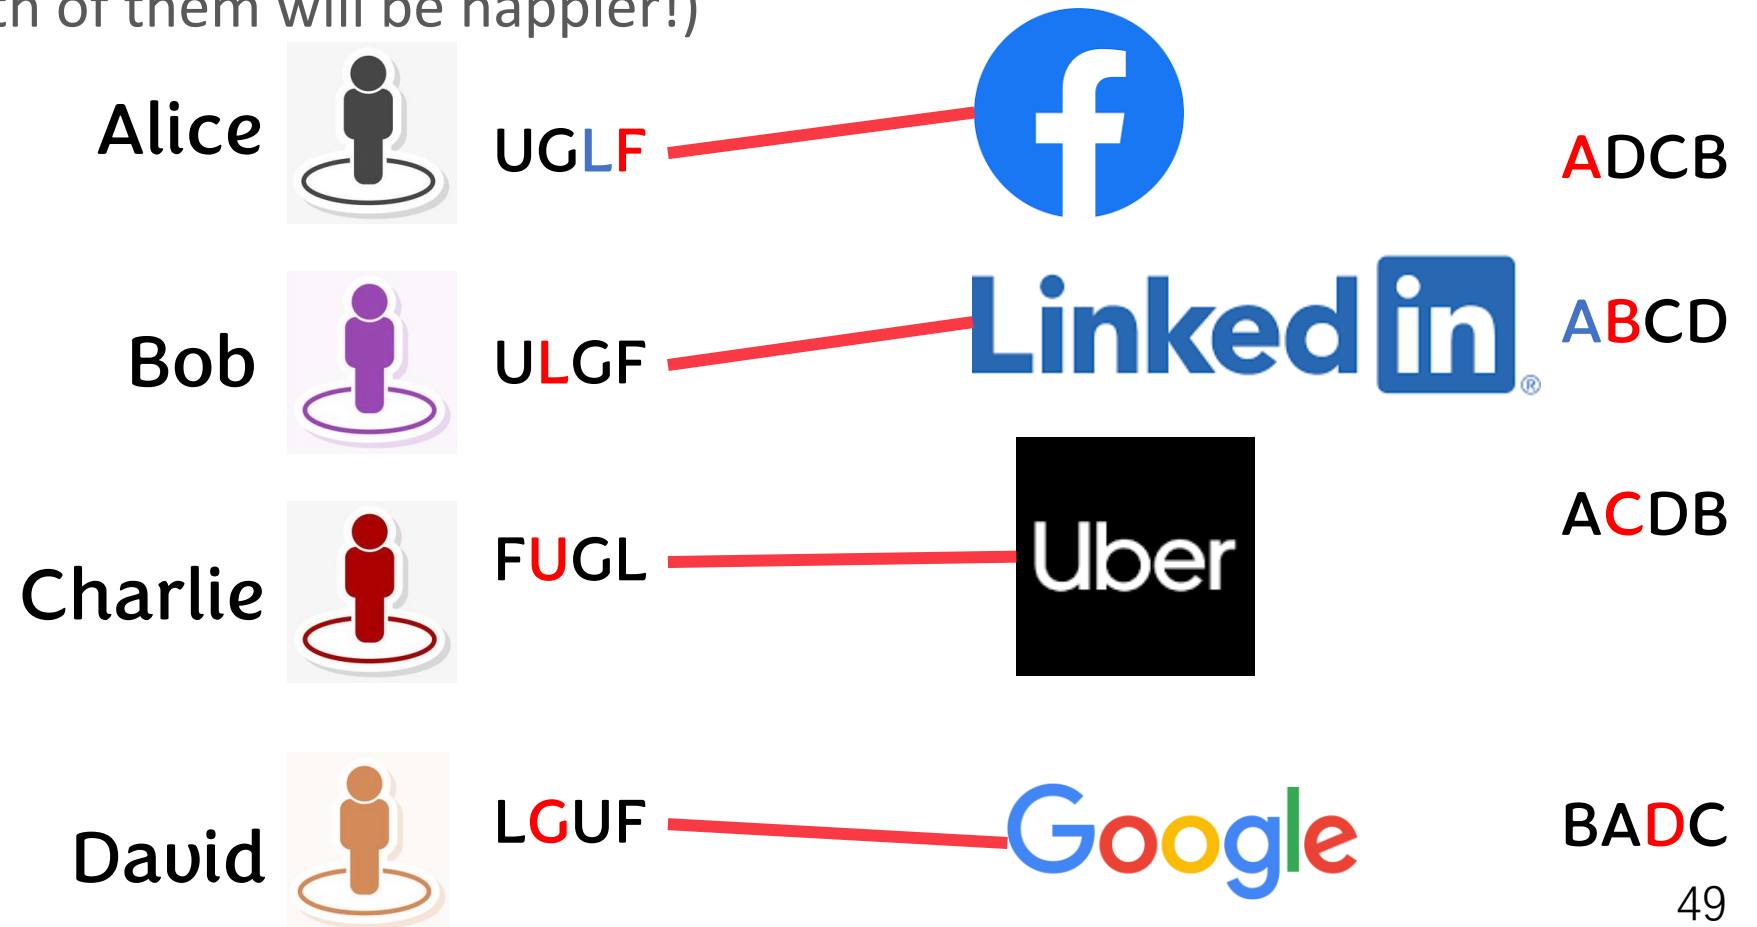
Finding a good match

- A stable match means that, there is no matching of:

- $A \Leftrightarrow 1, B \Leftrightarrow 2$, where A prefers 2 better than 1, and 2 prefers A better than B
- (if A move to 2, both of them will be happier!)

An unstable matching:

- Alice prefers LinkedIn better than Facebook
- LinkedIn prefers Alice better than Bob (current employee)
- So Alice should move to LinkedIn, and LinkedIn should also accept Alice (replace Bob)




Finding a good match

- A stable match means that, there is no matching of:

- $A \Leftrightarrow 1, B \Leftrightarrow 2$, where A prefers 2 better than 1, and 2 prefers A better than B
- (if A move to 2, both of them will be happier!)

A stable matching:

- No one has a better choice
- E.g., Bob will work at LinkedIn, he/she will be happier to be at Uber
- However, Uber got its best choice, so it will not accept Bob
- No one will move



Finding a good match (Gale-Shapley algorithm)





- **Round 1:**

- All candidates go to the company on the top of its list for an interview
- For any company:
 - No candidate comes: do nothing
 - 1 candidate: accept
 - Multiple candidate: choose the one based on its preference list. Reject all others.
- Candidate: if got rejected, remove the company from the list.

- **Round 2 - ?**

- All candidates goes to the company on the top of its list for an interview
- For any company:
 - No candidate comes: do nothing
 - Given any candidate: look at the one x with the highest preference. If x is better than the current employee (or no current), replace (yes, cold-blooded employer!). Otherwise, keep the current.
- Candidate: if got fired or rejected, remove the company from the list.

Finding a good match

Alice		UGLF
Bob		ULGF
Charlie		FUGL
David		LGUF



ADCB



ABCD




ACDB




BADC


Finding a good match




Finding a good match




Finding a good match




Finding a good match



Finding a good match



Finding a good match



Gale-Shapley Algorithm

David



LCUF

- The algorithm must give you a stable matching

First, it must stop

- In each round, either 1) all applicants have a job, which means the algorithm stops, or 2) at least one applicant will apply for a job
- We will show that, 2) can't continue forever
 - An applicant will traverse all companies in the list, so one day s/he should have done interviews in all companies
 - For a company, as long as anyone applied for the job, the position will be taken
 - This means, when any applicant finish the list, all companies must have their employees hired
 - Contradiction!

Gale-Shapley Algorithm

Alice
Linkedin>Facebook



Second, it must be stable!

Bob

• A company's employee is only getting better over time
• An applicant's job is only getting worse over time
• Assume \exists Alice at Facebook and Bob at LinkedIn, where Alice likes LinkedIn more than Facebook, and LinkedIn likes Alice more than Bob
• So Alice must have applied for job at LinkedIn before Facebook
• If LinkedIn didn't make an offer to Alice, it must have a better choice than Alice

Alice>Bob

- Note that
 - A company's employee is only getting better over time
 - An applicant's job is only getting worse over time
- Assume \exists Alice at Facebook and Bob at LinkedIn, where Alice likes LinkedIn more than Facebook, and LinkedIn likes Alice more than Bob
 - So Alice must have applied for job at LinkedIn before Facebook
 - If LinkedIn didn't make an offer to Alice, it must have a better choice than Alice
 - So LinkedIn couldn't end up offering the job to someone worse than Alice (Bob)
 - Contradiction!

Stable Matching Problem (Stable Marriage Problem)

- This algorithm is usually described for matching a set of man and women, in order to find “stable marriage”
 - There doesn't exist Alice and Bob, such that Alice like Bob more than her husband, and Bob like Alice more than his wife (otherwise?)
- **Gale–Shapley algorithm gives a solution to such problems**

Stable Matching Problem (Stable Marriage Problem)

- So how does this model “marriage”?
 - In the first round,
 - each single man proposes to the woman he prefers most, and
 - each woman accept the suitor she prefers most and say "no" to all other suitors.
 - “accept” mean “provisionally engaged” or “in relationship” – not marrying!
 - In later rounds,
 - each single man proposes to the most-preferred woman who hasn’t said “no” to him
 - each woman accept the most preferred suiter if she was single, or if her most preferred suitor is better than her current partner, update
 - Finally the process stops, that would be a stable marriage!

Stable Matching Problem (Stable Marriage Problem)

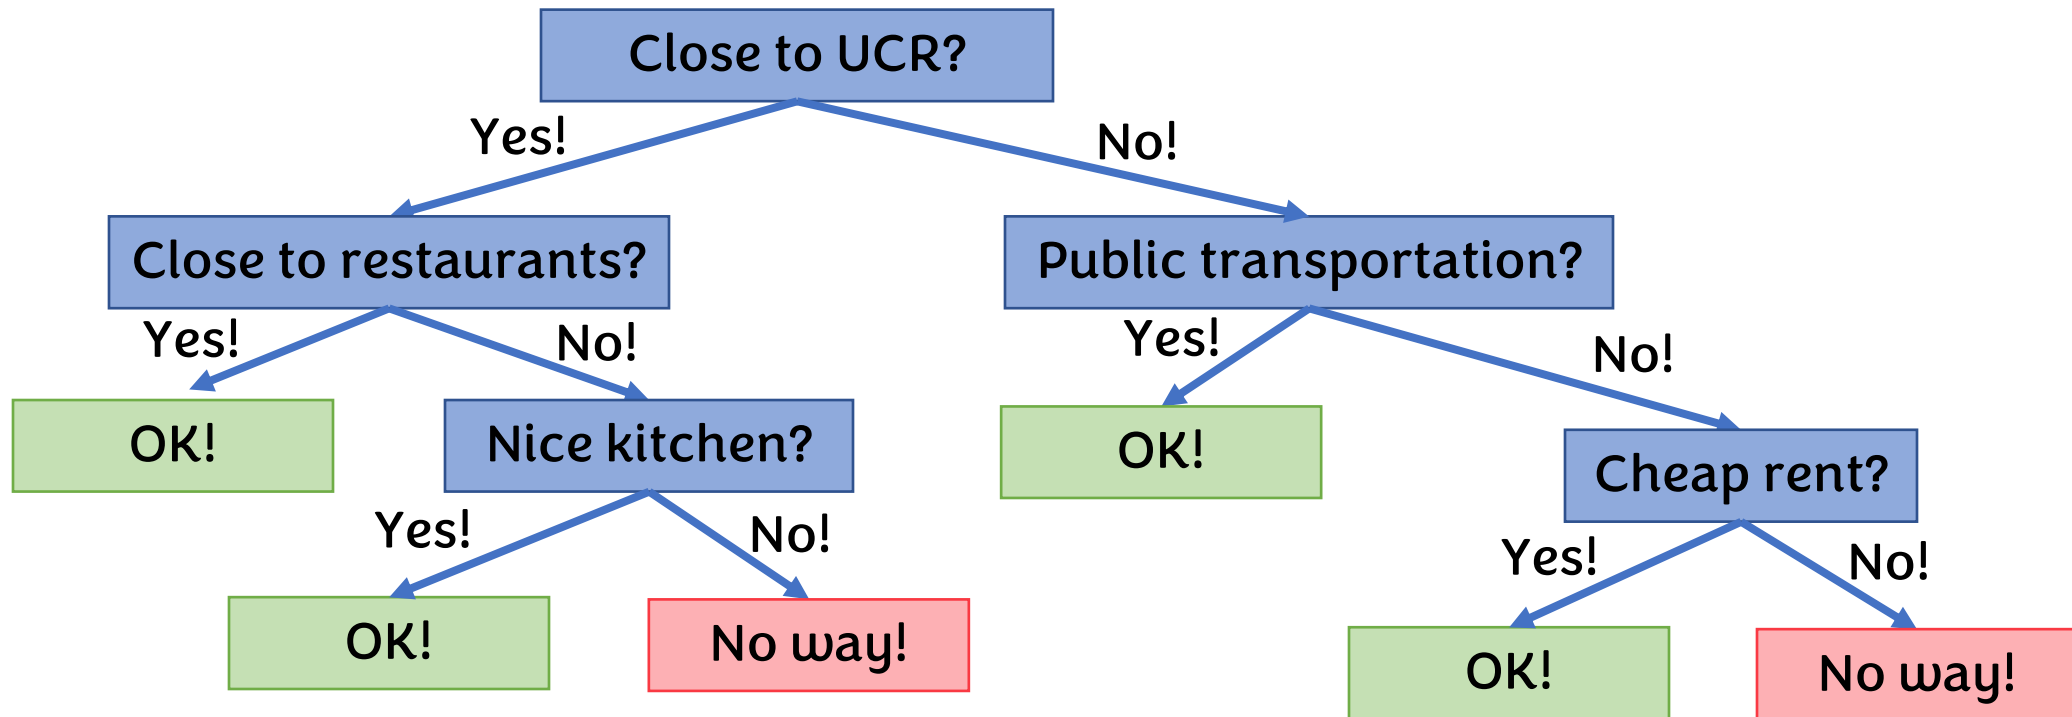
- **So how does this model “marriage”?**
 - It somehow indicates the following result for the algorithm
 - A man’s partner will get worse and worse (but they have the chance to propose to their top choice)
 - A woman’s partner will get better and better (but they can only choose among those that have proposed to them)
- **The algorithm is used in applicant-job matching (doctor-hospital) much earlier than it’s formally studied!**
 - A simple algorithm and is easy to implement
 - Somehow people feel like it will work intuitively
 - but don’t know why for a while

Stable Matching Problem (Stable Marriage Problem)

- This algorithm is usually described for matching a set of man and women, in order to find “stable marriage”
 - There doesn't exist Alice and Bob, such that Alice like Bob more than her husband, and Bob like Alice more than his wife (otherwise?)
- **Gale–Shapley algorithm gives a solution to such problems**
 - It does not need a centralized authority to “run” the algorithm
 - It only needs to tell each participant (or applicant/company) what the best strategy is. They will play the game automatically and get a stable solution!

Other interesting greedy algorithms

- **Decision tree: find the feature with the best “Information gain”**
 - The reduction in entropy or surprise by transforming a dataset
 - Intuitively: find a feature such that: one branch is mostly “yes” and the other is mostly “no”



Summary for today's lecture

- How to handle NP-hard problems or other problems with their best solutions being very slow?
- Approximation algorithms can be useful in this case, and oftentimes we can prove the quality of the solutions
- Sometimes it is also theoretically interesting to understand the impossible results for approximation
 - Take CS 215 and 219 (Spring 2023) for more details

Next lecture...

- **Data structure I: tournament trees and augmented trees**