


CS218: design and analysis of algorithms

Data Structures

Yan Gu


What data structures have we learned before?

- **Linked list**



- **Array**

- 1D, 2D, ...



- **Search tree**




- Binary search tree/multiway search tree/balanced search tree/AVL tree/red-black tree

- **Heap**

- Binary heap, Fibonacci heap, ...

- **Hash Table**

- Open addressing, close addressing, ...




Construction of Huffman Tree

- **Huffman(C)**
 - $n = |C|$
 - $Q = C$ // construct a **priority queue** of all character's frequency
 - for $i = 1$ to $n-1$
 - allocate a new node z
 - $z.\text{left} = x = \text{Extract-Min}(Q)$
 - $z.\text{right} = y = \text{Extract-Min}(Q)$
 - $z.\text{freq} = x.\text{freq} + y.\text{freq}$
 - **Insert**(Q, z)
 - return **Extract-Min**(Q) // **Root of the tree**

What priority queue will you use?

Construct Huffman tree using priority queues

Operation needed	Function name
Construct a priority queue (initialize) with n elements	construct (array A)
Find the smallest element and delete it	extract_min()
Insert an element	insert(x)



This is “**priority queue**”, it is an “**abstract data type**” (ADT).
It specifies an interface of functions

- In Huffman tree construction, we need a “**priority queue**”
- we call **construct** once
- **extract_min** and **insert** n times

Construct Huffman tree using priority queues

To implement a “**priority queue**”, we need some “**data structures**”.
Here are some possible implementations.

Data Structure	construct (array A)	extract_min()	insert(x)
Binary heap	Construct a heap from an array (heapify)	Read the root and delete it	Insert x into the heap
	$O(n)$	$O(\log n)$	$O(\log n)$
Balanced binary tree (e.g., AVL tree)	Construct a tree from an array	Chase the left-most branch to find min and delete it	Insert into the tree
	$O(n \log n)$	$O(\log n)$	$O(\log n)$
Sorted array	Sort the array	Find the first element in the array and mark it deleted	Insert x into the middle of the array to keep the order
	$O(n \log n)$	$O(1)$	$O(n)$
Unsorted array	Put everything in the array (nothing to do)	Traverse the array to find the min and mark it deleted	Put x at the end
	$O(1)$	$O(n)$	$O(1)$


Is the abstraction of “priority queue” useful?

- Is it only used in Huffman tree construction?
- Do you know other algorithms using priority queues?
- Indeed, a lot of greedy algorithms essentially use priority queues

Dijkstra's algorithm and Prim's algorithm

- Given a graph and a source vertex s , find the shortest distance from s and all the other vertices

- $\delta(u) = \infty$ for all $u \in V$, 0 for $\delta(s)$
- $S = \emptyset$
- $Q = \{s\}$
- **while** $Q \neq \emptyset$
 - $u = \text{Extract-Min}(Q)$
 - $S = S \cup \{u\}$
 - **for each** $v \in N(u)$
 - $\delta(v) = \min \{\delta(v), \delta(u) + w(s, v)\}$
 - **Insert/Update** $\delta(v)$ in Q



- $\delta(u)$: tentative distance
- S : settled set
- Q : priority queue
- $w(u, v)$: weight of edge from u to v
- $N(u)$: neighbor set of u

Is the abstraction of “priority queue” useful?

- **A lot of greedy algorithms can essentially use priority queues**
 - Because they make “greedy” choices
- **Candy buying: find the cheapest candy**
 - construct, extract_min. No need to insert or update
 - Using sorted array, heap, or balanced binary tree will all work ($O(n \log n)$ cost)
- **Dijkstra’s and Prim’s algorithm**
 - extract_min, insert, update. No need to construct (start from a single element)
 - Using sorted array gives you $O(n^2)$ time.
 - Using binary heap or balanced binary tree gives you $O(m \log n)$ time
 - Using something called “Fibonacci heap” gives you $O(n \log n + m)$ time
- **Huffman Code: Construct, extract_min, insert**
- **Optimization with submodularity: CELF (lazy update)**

What other abstract data types do you know?

Example: ADTs in STL:

- **Ordered set/maps**
 - (`std::map`, `std::set`)
 - implemented by red-black trees
- **Unordered set/maps**
 - (`std::unordered_map`, `std::unordered_set`)
 - implemented by hash tables

Why they are called “`std::map` / `std::unordered_map`” instead of “`std::red_black_tree` / `std::hash_table`”?

What’s the difference between `sets` and `maps`?

Use ADT and data structures for algorithm design

- Your algorithm may need to access and organize data:
 - How to store data?
 - What query to support? (lookup, findMin, findSum, ...)
 - What update to support? (insertion, deletion, filter, multi_insert, delete_min, ...)
- Based on the functions needed, you define an **abstract data type (ADT)**!
 - You don't care how they are supported (data organization/algorithms/cost bounds) for an ADT
- Then you find a **data structure** to support them
 - A concrete way to organize data and concrete algorithms to support the functions
 - They may have different costs for different functions
 - So based on your algorithm, you choose the best data structure

Examples of ADT

- FIFO Queue
 - Deque (double-ended queue)
 - Stack
 - Priority queue
 - Ordered set/map
 - Unordered set/map
-
- (sometimes we say a queue or a stack is a data structure when we are talking about a concrete implementation)

In this course

- **Winning trees**
 - An implementation of priority queue
 - Easy to implement, same bound as binary heap
- **Augmented trees**
 - Range-related queries: 1D range max/min/sum
 - Rank/selection on trees

Winning Tree


Priority queue

- **Store a set of keys**
 - find/delete smallest/largest key
 - update the keys
 - Insert/delete

Operation needed	Function name
Construct a priority queue (initialize) with n elements	construct (array A)
Find the smallest element and delete it	extract_min()
Insert an element	insert(x)


Binary Heap

- Organize all keys in a complete binary tree
- The key at node x is smaller than its children
- Once updated, an element needs to move upwards or downwards (**heapify**)
- Not easy to implement
- Min/max query applicable only to the entire set
 - Cannot support min/max of a range




Winning tree

- Store all elements at the leaves
- Each internal node is a **competition**
 - the one wins (the smaller one) will be recorded at the node
- To make it easier to be stored in an array, add some dummy nodes to make the size 2^k for integer k



Winning tree


- Store all elements at the leaves
- Each internal node is a competition
 - the one wins (the smaller one) will be recorded at the node
- To make it easier to be stored in an array, add some dummy nodes



- **Insertion**
 - add at the end
 - re-compute all its ancestors
- **Deletion**
 - mark it as ∞
 - re-compute all its ancestors
- **Update**
 - update its key
 - re-compute all its ancestors


Winning tree - construct

- $O(n)$ time




Winning tree - update

- $O(\log n)$ time




Winning tree: insertion

- $O(\log n)$ time




Winning tree - deletion

- $O(\log n)$ time
- If we want to make it more space-efficient, we can move the last element back to the slot of the deleted element




Winning tree



- Can also use an array to store it (similar to binary heap)
- Two children of $A[i]$: $A[2i]$ and $A[2i+1]$ (can also start from 0)

Index	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Value	2	2	8	2	6	8	∞	2	4	9	6	8	10	∞	∞

Winning tree



- Can also use an array to store it
- Two children of $A[i]$: $A[2i]$ and $A[2i+1]$ (can also start from 0)
- We don't need dummy nodes!! Use a **complete binary tree**

Index	1	2	3	4	5	6	7	8	9	10	11
Value	2	6	2	6	8	2	4	9	6	8	10


Winning tree vs. binary heap

- Same asymptotical bound for insertion/deletion/update/extract_min
- **Winning tree takes more space, and is slightly slower in practice**
 - $O(\log n)$ cost is tight for winning trees
- **Winning tree is much simpler to implement**
 - Essentially, it only needs one operation
- **However, winning tree is more general than binary heap**

Winning tree – range query

- Find the min of the 2nd – 6th elements?


min(4, 9, 6, 8, 10)



Winning tree – range query

- Find the min of the 2nd – 6th elements?


min(4, 9, 6, 8, 10)



Winning tree – range query

- Find the min of the 2nd – 6th elements?


min(4, 6, 8)



Winning tree – range query

- Find the min of the first 6 elements?


$\min(2, 4, 9, 6, 8, 10)$



Winning tree – range query

- Find the min of the first 6 elements?


$\min(2, 4, 9, 6, 8, 10)$



Winning tree – range query

- Find the min of the first 6 elements?

$\text{min}(2, 8)$



Winning tree vs. binary heap

- Same asymptotical bound for insertion/deletion/update/extract_min
- **Winning tree takes more space, and is slightly slower in practice**
 - $O(\log n)$ cost is tight for winning trees
- **Winning tree is much simpler to implement**
 - Essentially, it only needs one operation

Augmented search trees

Range sum query


- Given a set of key-value pairs, a query asks for the sum of values in between a key range

		25,2	16,10	12,11	15,17	7,7	9,2				
		8,9	4,4		17,3	13,9	10,2				
Key	4	7	8	9	10	12	13	15	16	17	25
value	4	7	9	2	2	11	9	17	10	3	2
Prefix sum	4	11	20	22	24	35	44	61	71	74	76


- $\text{range_sum}(7, 16) = 71 - 4 = 67$
- $O(\log n)$ time

Range max query


- Given a set of key-value pairs, a query asks for the max of values in between a key range
 - Cannot use subtraction!
- The set of key-value pairs can be updated
 - Insert/delete new key-values or update the values
- Search tree
 - Insertion/deletion/update in $O(\log n)$ time



Why we need search trees?



A B-tree structure maintaining ordering on keys.



A binary search tree structure maintaining ordering on keys.

- Organizing a set of data
- What is the benefit/disadvantage of using trees compared to arrays?
- What is the benefit/disadvantage of using trees compared to hash table?

Search Trees

- **Ordered!**

- The in-order traversal is sorted w.r.t. keys
- (cannot be achieved if we use hash tables)

- **Dynamic!**

- Insertion/deletion in $O(\log n)$ time
- (cannot be achieved if we use ordered arrays)

- **Efficient!**


- Insertion/deletion/lookup in $O(\log n)$ time
- Traversal in $O(n)$ time
- Get some info from / operate on the root of subtrees (do not need to touch all tree nodes)

- **However?**

- More space than flat arrays
- Worse locality since tree nodes are scattered

Balanced Binary Trees

- Binary: each tree node has at most **two** children
- Balanced: the tree has **bounded height**
 - Usually $O(\log n)$ for size n




hyphaene compressa

A **wild**
balanced
binary tree

Balanced Binary Trees


- Binary: each tree node has at most **two** children
- Balanced: the tree has **bounded height**
 - Usually $O(\log n)$ for size n



A **wild**
balanced
binary tree

Balanced Binary Trees


- Binary: each tree node has at most **two** children
- Balanced: the tree has **bounded height**
 - Usually $O(\log n)$ for size n




An **abstract**
balanced
binary tree

Binary Search Trees



- Lookup/insertion just follows the path
- Deletion may be more involved
- Need to rebalance the tree using rotation!



Tree rotation



Source: wikipedia




Source: https://www.researchgate.net/figure/A-single-left-rotation-and-a-double-left-rotation-a-b-and-c-are-elements-x-y-y0_fig4_220676747

AVL trees

- Invariant: for any node in the tree, the heights of its two subtrees differ by at most 1
- Height is bounded by $O(\log n)$

- $|h(T_L) - h(T_R)| \leq 1$
 - $h(\cdot)$ denotes the height of the tree




Red-black trees

- **Invariants:**

- Red node only have black children;
- External nodes (leaves) are black;
- The black height from the root to any leaf is the same.

- **Tree height at most $O(\log n)$**



Augmented tree

Augmented data structures

- Very rarely we need to design a brand-new data structure
- We could augment some of the existing ones!
 - Directly use existing operations/analysis
 - With minimal additional information to be maintained


Range sum query

- Given a set of key-value pairs, a query asks for the sum of values in between a key range

Key	4	7	8	9	10	12	13	15	16	17	25
value	4	7	9	2	2	11	9	17	10	3	2

- Search tree**


- Insertion/deletion/update in $O(\log n)$ time



First of all, how to find all records in a given key range on a binary search tree?


Range query algorithm

- Report all entries in key range $[k_L, k_R]$.



```

range( $t, k_L, k_R$ ) {
   $r = t.root$ ; if (! $r$ ) return;
  if ( $k_R < r$ ) return range( $t.left, k_L, k_R$ ); // total range in the left branch
  if ( $k_L > r$ ) return range( $t.right, k_L, k_R$ ); // total range in the right branch
  add rangeR( $t.left, k_L$ ) to result
  add  $t.root$  to result
  add rangeL( $t.right, k_R$ ) to result
}
    
```





rangeL and rangeR functions

- rangeL: everything $\leq k$

```
rangeL(t, k) {  
  r = t.root; if (!r) return 0;  
  if (k < r.key) return rangeL(t.left, k); // total range in the left branch  
  if (k ≥ r.key) {  
    add t.left to result //the whole left subtree should be included  
    add t.root to result  
    if (k > r.key) add rangeL(t.right, k) to result  
  }  
}
```

Totally on the left

The left subtree is
totally included



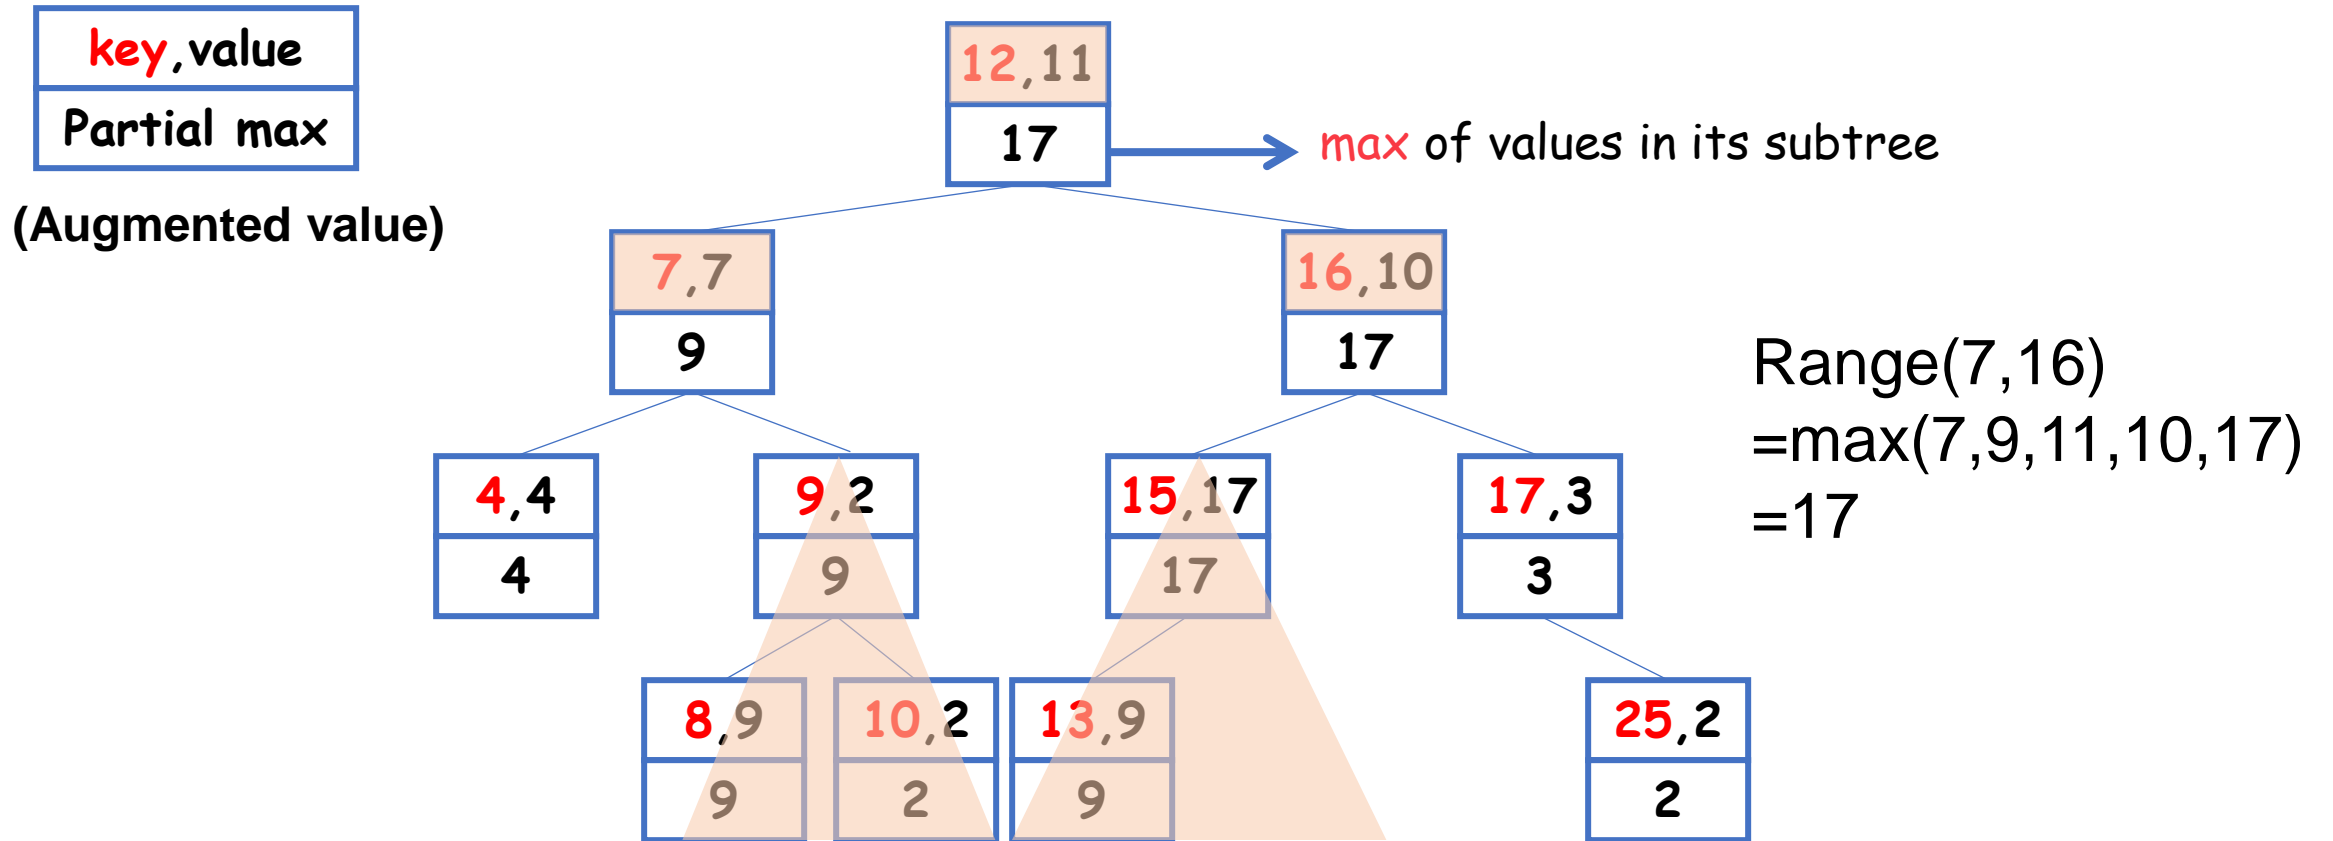
Range query: time complexity

- $O(\log n)$ time to find the relevant subtrees
- $O(k)$ time to traverse and output (k is the output size)
- Total cost: $O(\log n + k)$

**What if we only need the
sum/max/count of the range?**

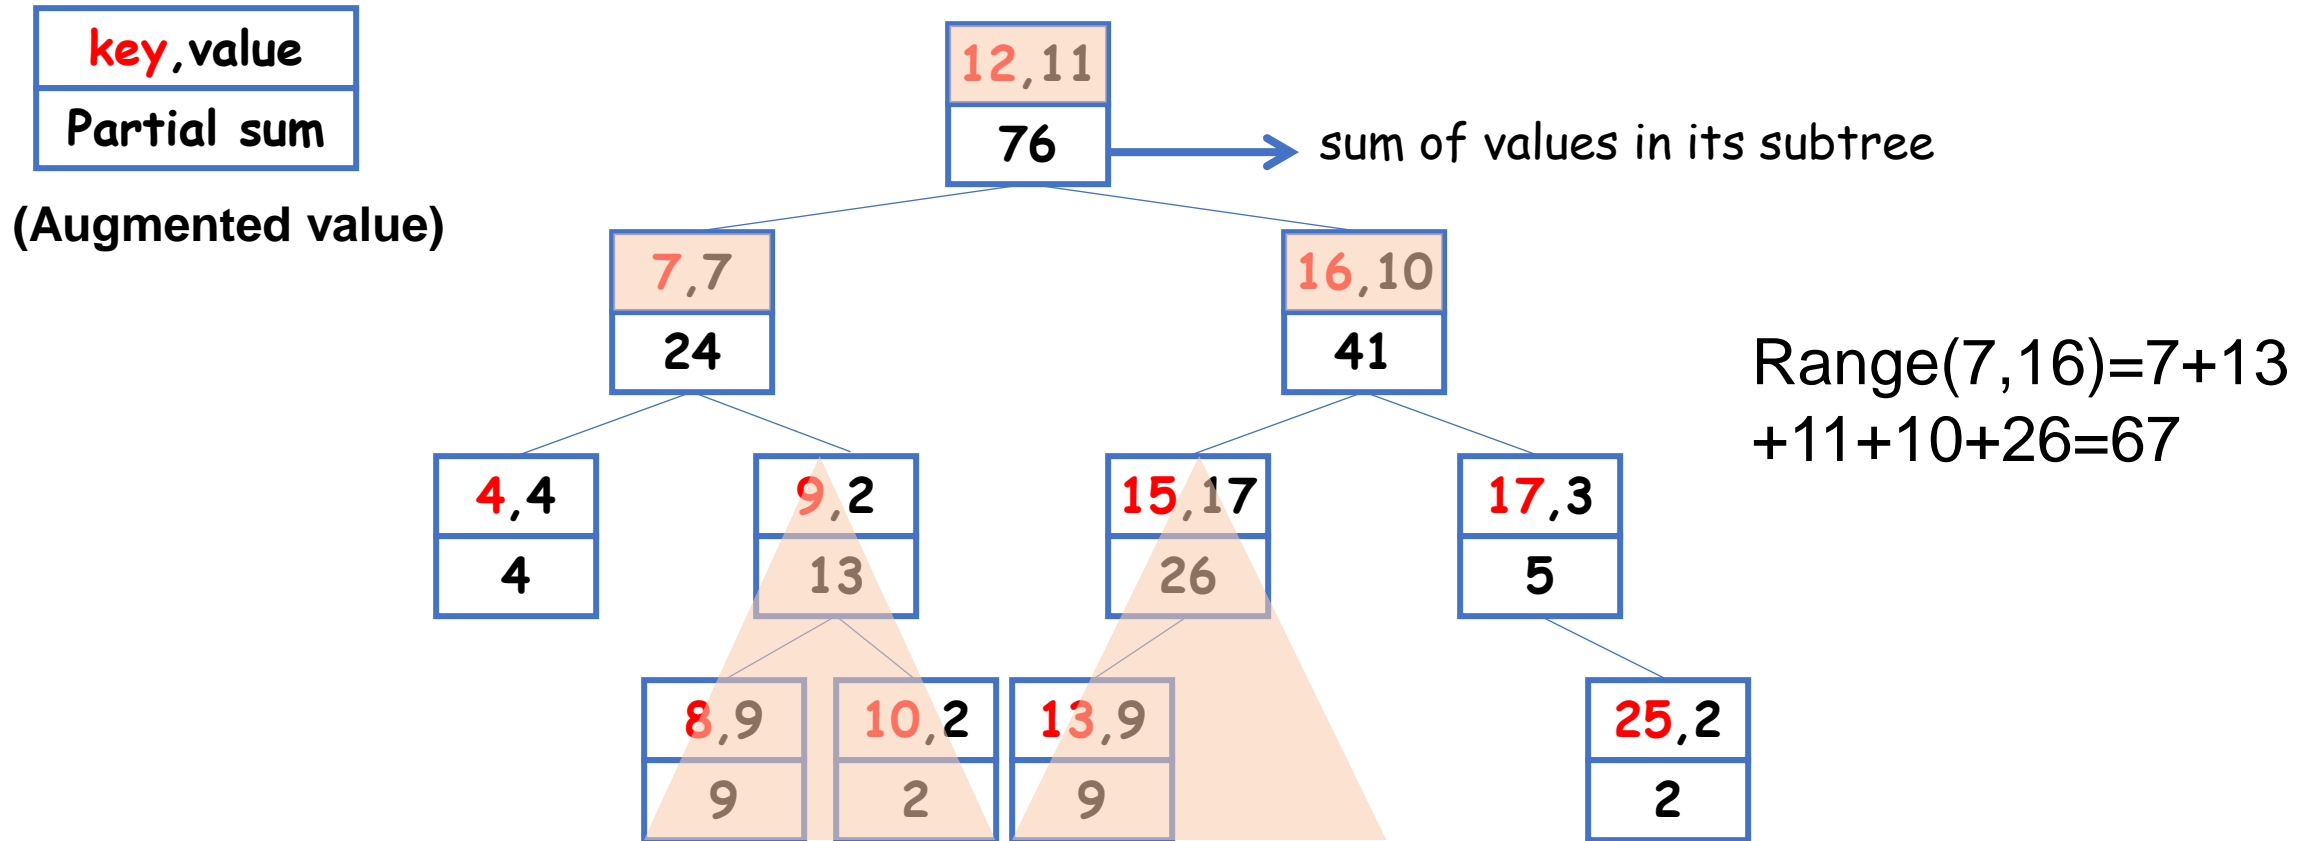
Augmented Trees for Range Max

- Storing the max in each tree node for answering range sum queries




Augmented Trees for Range Sum

- Storing the sum in each tree node for answering range sum queries



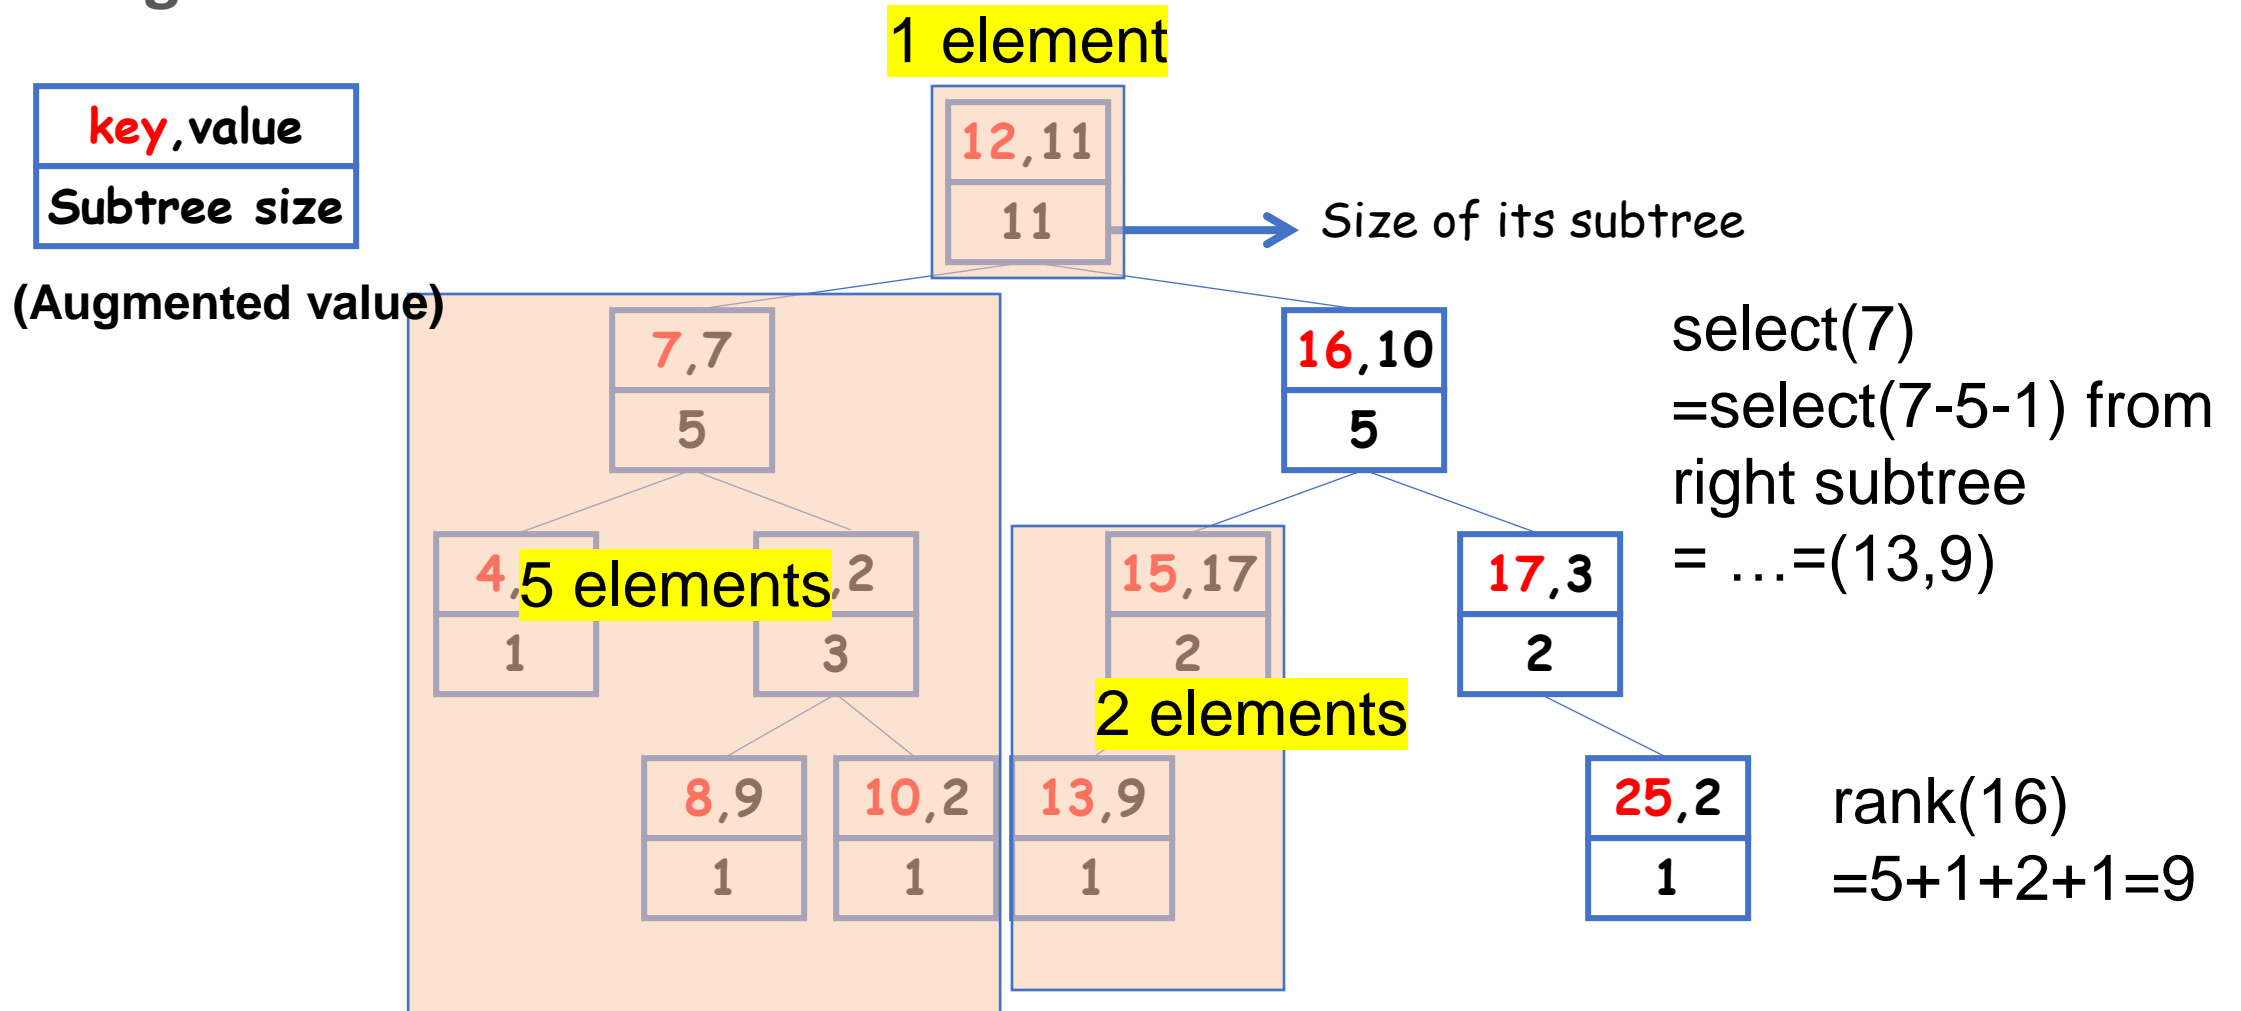
Range **max** query algorithm

- When the whole subtree is included in the result, sometimes we don't need to traverse the tree – use the **augmented value**!
- E.g., if we want the max of values
 - Store at each node the max value
 - When we say “add the subtree to the result”, directly read the max value (augmented value) from the subtree root, and add the value to the result
 - Takes only $O(\log n)$ time!



Augmented Trees for rank report / select k-th element

- Storing the subtree size



Maintain augmented values

- In any construction/insertion/deletion/update ... algorithms
- Any modification will update all related tree nodes on the path
- Asymptotically the same time bound
- **aug_val** of a node can be computed by its two children and the node:
 - $\text{min} = \min(\text{leftmin}, \text{rightmin}, \text{rootvalue})$
 - $\text{sum} = \text{leftsum} + \text{rightsum} + \text{rootvalue}$
 - $\text{size} = \text{leftsize} + \text{rightsize} + 1$

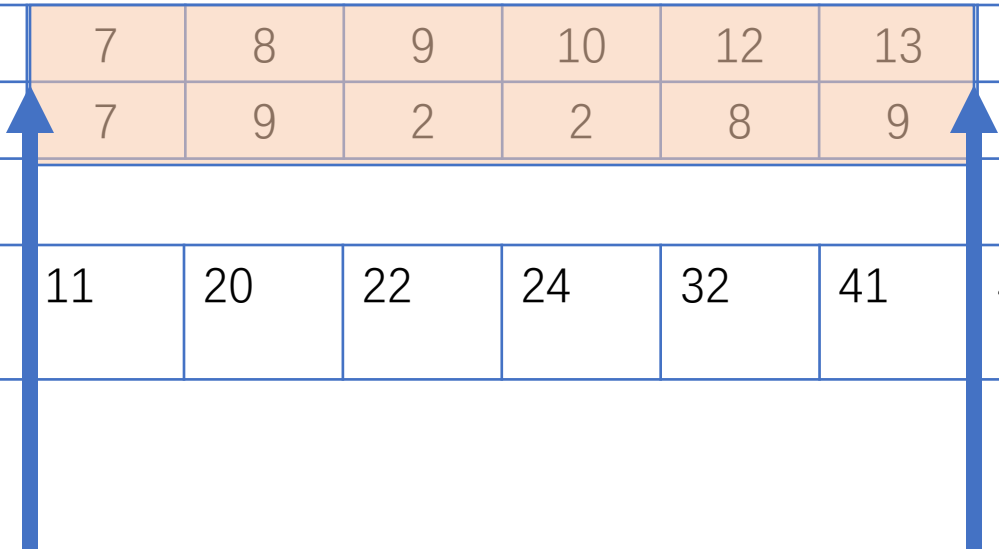
Augmented tree for range queries

- Range sum/min/max/... queries
- Can be combined with any searched tree data structure
 - Base data structure can be AVL, red-black tree, etc.
- Using different augmentations we can get different functionalities

Revisit of the range sum/max/rank query

- Given a set of key-value pairs, a query asks for the sum of values in between a key range

Key	4	7	8	9	10	12	13	15	16	17	25
value	4	7	9	2	2	8	9	7	10	3	2
Prefix sum	4	11	20	22	24	32	41	48	58	61	63

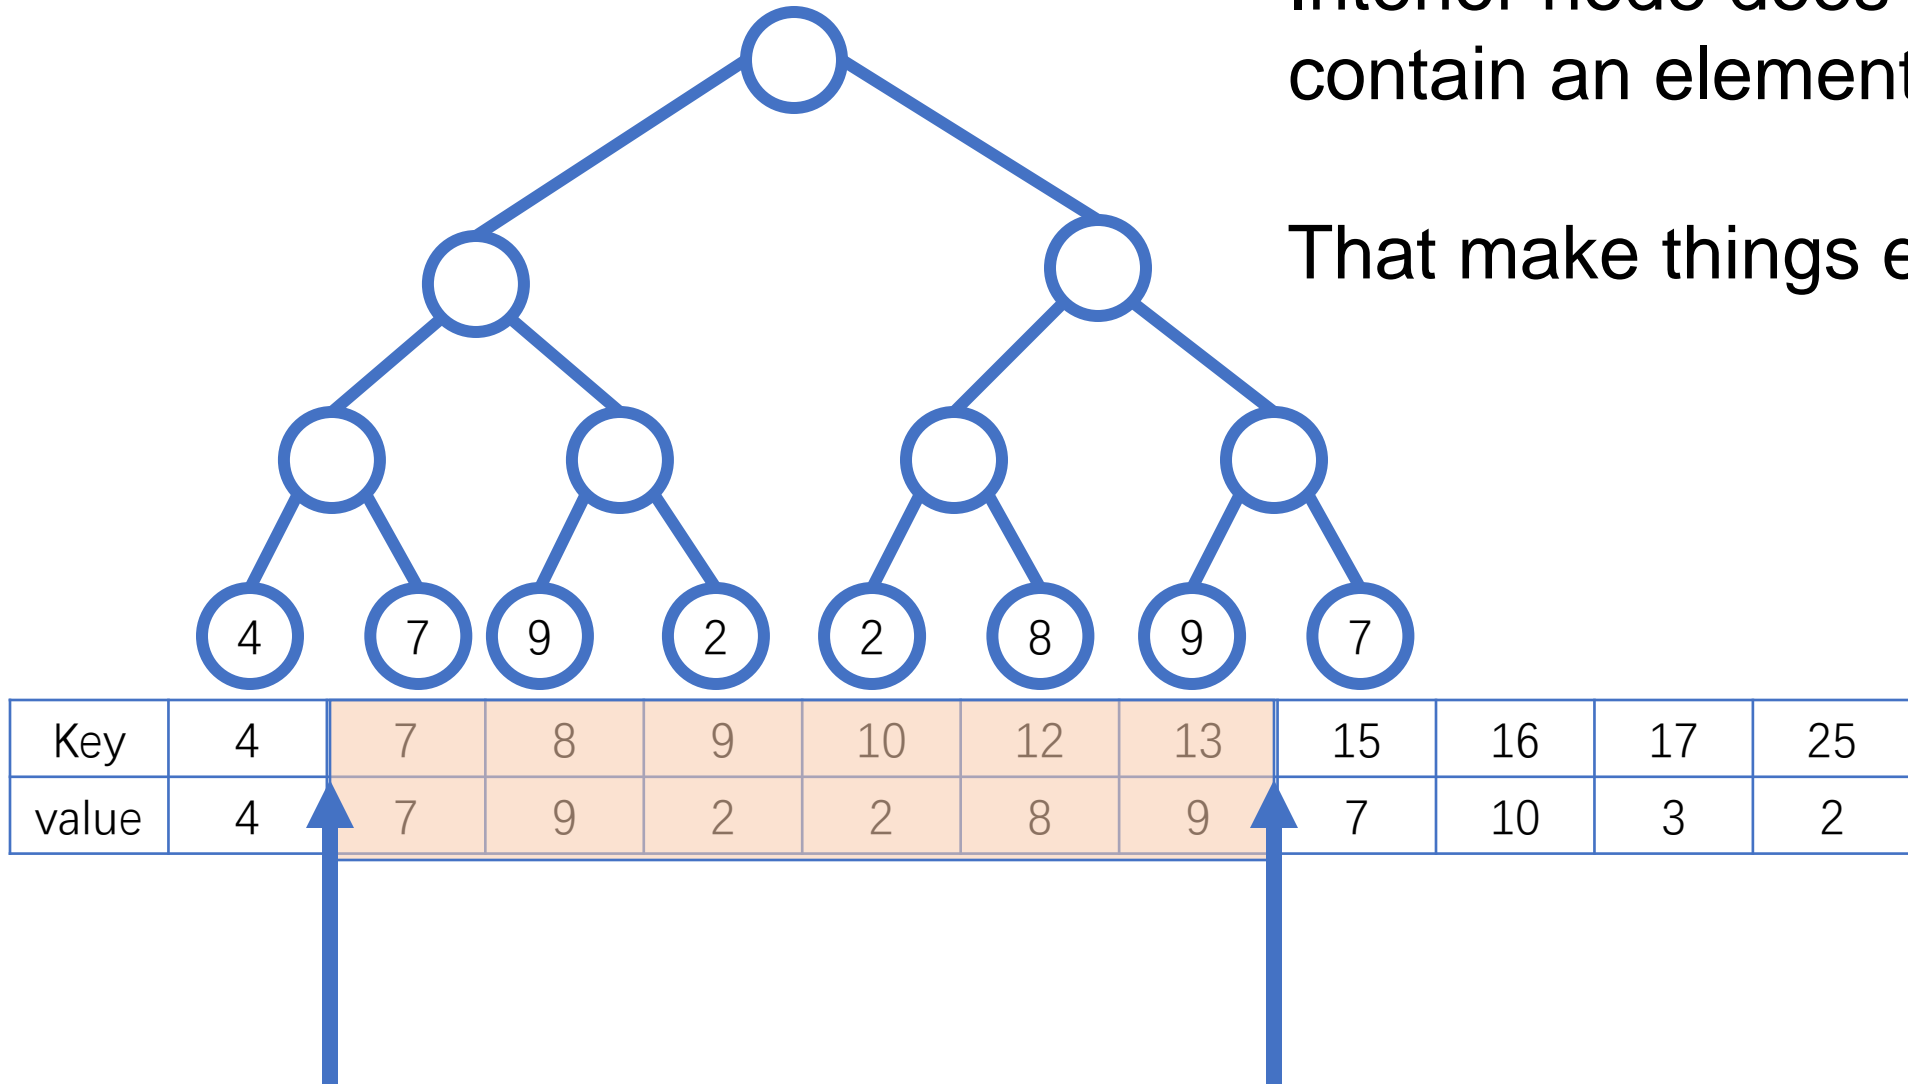


- $\text{range_sum}(7, 13) = 41 - 4 = 37$
- $O(\log n)$ time


Build a winning tree for the query!

Interior node does not
contain an element!

That make things easier!



Range query on a winning tree



```
rangeMax( $t, k_L, k_R$ ) {  
   $r = t.\text{root}$ ; if ( $r$  is null) return;  
  if ( $t.\text{min} < k_R$  and  $t.\text{max} < k_L$ ): return;  
  if ( $k_L \leq t.\text{min}$  and  $t.\text{max} \leq k_R$ ):  
     $\text{ans} = \max(\text{ans}, t.\text{augval})$  and return;  
  if ( $k_L \leq r \leq k_R$ ):  $\text{ans} = \max(\text{ans}, r)$ ;  
  rangeMax( $t.\text{left}, k_L, k_R$ );  
  rangeMax( $t.\text{right}, k_L, k_R$ );  
}
```

Winning tree itself is a special augmented tree

- It can be used directly as a priority queue because the extract-Min for priority queue is a special range query
- We can augment the winning tree differently for different queries; we can augment multiple fields simultaneously
- The left-to-right order of the leaves in a winning tree can be used to represent a list, and augmenting the winning tree can solve all list query problems
- If we pre-sort all elements or the key range is fixed (e.g., $[1, \dots, n]$), winning tree can be used as a static search tree (**discretization**)

Summary

Winning tree + augmentation is almost sufficient for any data structure to maintain 1D data

- **Winning tree stores the elements in all leaves in a complete binary tree**
 - Unordered: binary heap
 - A list: list queries (static)
 - A sorted list: static binary search trees
- **Augmentation supports range and rank queries (no matter what tree that is)**
 - For the whole list: stored in the root
 - For a range (k_L, k_R) : use the algorithm that visit $O(\log n)$ subtrees
 - Supports min/max/sum/rank/...
 - For list-all, the cost is $O(\log n + k)$ where k is the output size

The next lectures ...

- **Dynamic programming!**
- **(For the next four lectures)**
- **Training part of Homework 2 due tomorrow**